

# Advanced Programming

## Introduction to Erlang

Oleks Shturmov  
oleks@oleks.info

Slides adapted from Ken Friis Larsen  
kflarsen@diku.dk

Department of Computer Science  
University of Copenhagen

September, 2021

# Before we begin

- ▶ **Windows users**

- ▶ Add C:\Program Files\erl-24.0\bin to your PATH environment variable

- ▶ **Everyone**

dialyzer --build\_plt --apps erts kernel stdlib \  
crypto compiler

- ▶ Dialyzer is a tool for type-checking Erlang code
- ▶ This may take up to 20 minutes, exactly enough

# Part I

## Introduction

# The Erlang Programming Language



- ▶ Developed in the early 1980s, by the guys above<sup>1</sup>
- ▶ All while working at Ericsson, programming telephone switches
- ▶ Useful for distributed, fault-tolerant systems, in general
- ▶ Open-sourced in 1998, still maintained by Ericsson
- ▶ Used today by WhatsApp (Facebook), Nintendo, Discord, etc.

---

<sup>1</sup>Image source: [https://www.youtube.com/watch?v=rYkI0\\_ixRDc](https://www.youtube.com/watch?v=rYkI0_ixRDc)

# Erlang Customer Declaration

Erlang is a:

- ▶ a concurrency-oriented language
- ▶ dynamically typed
- ▶ with a strict functional core language

# On the matter of Erlang syntax

- ▶ Syntax heavily inspired by the Prolog programming language
- ▶ **Semantically, Erlang bears little resemblance to Prolog**
  - ▶ Prolog is a *logic programming language*
  - ▶ A distinct programming paradigm, beyond the scope of this course
- ▶ We won't delve on this history, except to justify peculiar syntax

## Part II

### Basic Concepts

# erl

## On Windows, the Erlang executable

- ▶ Starts an Erlang runtime system (i.e., an Erlang node)
  - ▶ Distributed (multi-node) Erlang is beyond the scope of this course
- ▶ Drops you into an Erlang shell for that node

```
$ erl
```

```
Erlang/OTP 24 [erts-12.0.3] [source] [64-bit] [...]
```

```
Eshell V12.0.3 (abort with ^G)
```

```
1> 21+21.
```

```
42
```

```
2>
```

- ▶ Use one of the following options to quit/kill erl:
  - ▶ Enter the command `q()`.
  - ▶ Press `Ctrl+g`, and enter the command `q` (quit)
  - ▶ Press `Ctrl+c`, and enter the command `a` (abort)



# Fundamental Stuff

- We have integers, floating-point numbers, and arithmetic:

```
1> 21+21.
```

```
42
```

```
2> 3/4.
```

```
0.75
```

```
3> 5 div 2.
```

```
2
```

- We have lists:

```
4> [21,32,67] ++ [100,101,102].
```

```
[21,32,67,100,101,102]
```

- Strings are just lists of characters, and characters are just integers:

```
5> "Sur" ++ [112, 114, $i, $s, $e].
```

# Fundamental Stuff

- We have integers, floating-point numbers, and arithmetic:

```
1> 21+21.
```

```
42
```

```
2> 3/4.
```

```
0.75
```

```
3> 5 div 2.
```

```
2
```

- We have lists:

```
4> [21,32,67] ++ [100,101,102].
```

```
[21,32,67,100,101,102]
```

- Strings are just lists of characters, and characters are just integers:

```
5> "Sur" ++ [112, 114, $i, $s, $e].
```

```
"Surprise"
```

# Names (Variables)

- Names (variables) start with an upper-case letter

```
1> Homer = "Homer".
```

```
Homer
```

```
2> X=5, Y=2, X*Y.
```

```
10
```

- Once assigned, variables cannot be re-assigned<sup>2</sup>

```
3> X=3, Y=2, X*Y.
```

```
** exception error: no match of right hand side value 3
```

---

<sup>2</sup>Perhaps another relic of the Prolog past, or part of the “strict functional core”

# What to do if you mess up in erl?

- ▶ Eshell is quite forgiving – you can just let exceptions happen
- ▶ In non-exceptional cases<sup>3</sup>, you might be tempted to kill erl...
- ▶ A better option is to press Ctrl+g:

```
4> X/0
4> % Pressing Ctrl+g ...
User switch command
-> h
  c [nn]          - connect to job
  i [nn]          - interrupt job
  ...
  q              - quit erlang
  ? | h          - this message
-> i
-> c
** exception exit: killed
4> X/2.
2.5
```

- ▶ Bonus: you get to keep your names (variables)!

---

<sup>3</sup>Long-running, live-, or dead-locked commands

# Tuples and Atoms

- ▶ Erlang uses curly braces for tuples:

```
1> {"Bart", 9}.  
{"Bart",9}
```

- ▶ **Atoms** are used to represent non-numerical constant values (like enums in C and Java). Atom is a sequence of alphanumeric characters (including @ and \_) that starts with a lower-case letter (or is enclosed in single-quotes):

```
2> bloody_sunday_1972.  
bloody_sunday_1972  
3> [{bart@simpsons, "Bart", 9}, {'HOMER', "Homer", 42}].  
[{bart@simpsons,"Bart",9},{ 'HOMER', "Homer",42}]
```

# Patterns

- As in Haskell, we can use patterns to take things apart:

```
1> P = {point, 10, 42}.
```

```
2> [ C1, C2, C3 | Tail ] = "Homer".
```

```
"Homer"
```

```
3> C2.
```

```
111
```

```
4> Tail.
```

```
"er"
```

```
5> {point, X, Y} = P.
```

```
{point,10,42}
```

```
6> X.
```

```
10
```

```
7> Y.
```

```
42
```

# List Comprehensions

```
1> Digits = [0,1,2,3,4,5,6,7,8,9].
```

```
[0,1,2,3,4,5,6,7,8,9]
```

```
2> Evens = [ X || X <- Digits, X rem 2 == 0].
```

```
[0,2,4,6,8]
```

```
3> Cross = [{X,Y} || X <- [1,2,3,4], Y <- [11,22,33,44]].
```

```
[{1,11}, {1,22}, {1,33}, {1,44},
```

```
 {2,11}, {2,22}, ... ]
```

```
4> EvenXs = [{X,Y} || {X,Y} <- Cross, X rem 2 == 0].
```

```
[{2,11},{2,22},{2,33},{2,44},{4,11},{4,22},{4,33},{4,44}]
```

# Maps

```
1> M = #{ name => "Ken", age => 45}.
#{age => 45, name => "Ken"}
2> ClunkyName = maps:get(name, M).
"Ken"
3> #{name := PatternName} = M.
4> PatternName.
"Ken"
5> #{name := Name, age := Age} = M.
#{age => 45, name => "Ken"}
6> {Name, Age}.
{"Ken", 45}
7> Wiser = M#{age := 46}.
#{age => 46, name => "Ken"}
8> WithPet = M#{pet => {cat, "Toffee"}}.
#{age => 46, name => "Ken", pet => {cat, "Toffee"}}
```



# Functions

- Remember the move function from Exercise Set 0 (Haskell)?

```
move :: Direction -> Pos -> Pos
```

```
move North (x,y) = (x, y+1)
```

```
move West  (x,y) = (x-1, y)
```

# Functions

- Remember the move function from Exercise Set 0 (Haskell)?

```
move :: Direction -> Pos -> Pos
```

```
move North (x,y) = (x, y+1)
```

```
move West  (x,y) = (x-1, y)
```

- In Erlang:

```
move(north, {X, Y}) -> {X, Y+1};
```

```
move(west, {X, Y}) -> {X-1, Y}.
```

(note that we use semicolon to separate clauses, and period to terminate a declaration).

# Functions

- Remember the move function from Exercise Set 0 (Haskell)?

```
move :: Direction -> Pos -> Pos
```

```
move North (x,y) = (x, y+1)
```

```
move West  (x,y) = (x-1, y)
```

- In Erlang:

```
move(north, {X, Y}) -> {X, Y+1};
```

```
move(west, {X, Y}) -> {X-1, Y}.
```

(note that we use semicolon to separate clauses, and period to terminate a declaration).

- Or naming a function literal:

```
Move = fun(Dir, {X,Y}) ->  
        case Dir of  
            north -> {X, Y+1};  
            west  -> {X-1, Y}  
        end  
    end.
```

# Modules

- ▶ If we want to declare functions (rather than naming literals) then we need to put them in a **module**.
- ▶ Modules are defined in `.erl` files, for example `somemodule.erl`:  

```
-module(warmup).  
-export([move/2]).
```

```
move(north, {X, Y}) -> {X, Y+1};  
move(west, {X, Y}) -> {X-1, Y}.
```
- ▶ Note, how we **specify the arity** of functions on export

## Compiling Modules

- Using the function `c`, we can compile and load modules in the Erlang shell:

```
1> c(warmup).  
{ok,warmup}
```

- We can now call functions from our module:

```
2> warmup:move(north, {0,0}).  
{0,1}
```

- Or use them with functions from the standard library:

```
3> Moves = [{north, {1,1}}, {west, {43,42}},  
            {north, {23,22}}].  
4> lists:map(fun({Dir,Pos}) ->  
              warmup:move(Dir, Pos) end,  
              Moves).  
[{1,2},{42,42},{23,23}]
```

## Part III

# Type Checking

# Erlang is dynamically typed

- ▶ The type of a value is determined at runtime
- ▶ Using values of wrong types induces runtime errors
- ▶ Perhaps, another relic of the Prolog past...
- ▶ However, static typing is principally complicated by Erlang's liberal communication primitives...

# Enter Dialyzer and TypEr

- ▶ Dialyzer and TypEr offer a **gradual type system** for Erlang
- ▶ Part of an Erlang code optimization effort
- ▶ The more types you know, the more can you optimize



# Enter Dialyzer and TypEr

- ▶ Dialyzer and TypEr offer a **gradual type system** for Erlang
- ▶ Part of an Erlang code optimization effort
- ▶ The more types you know, the more can you optimize
  
- ▶ Types form a (mathematical) lattice
- ▶ The top type is called `any()`
- ▶ All Erlang types conform to this type
- ▶ So we can safely give any value the type `any()`
- ▶ However, we can *gradually* increase the specificity of the types in our Erlang code as it evolves

## any()

- One straight-forward, albeit not very useful way to type move/2:

```
-module(warmup).
```

```
-export([move/2]).
```

```
-spec move(any(), any()) -> any().
```

```
move(north, {X, Y}) -> {X, Y+1};
```

```
move(west, {X, Y}) -> {X-1, Y}.
```

## any()

- One straight-forward, albeit not very useful way to type `move/2`:

```
-module(warmup).
```

```
-export([move/2]).
```

```
-spec move(any(), any()) -> any().
```

```
move(north, {X, Y}) -> {X, Y+1};
```

```
move(west, {X, Y}) -> {X-1, Y}.
```

- Let's see if Dialyzer would be happy:

```
$ erlc +debug_info warmup.erl
```

```
$ dialyzer warmup.beam
```

```
Checking whether the PLT ... is up-to-date... yes
```

```
Proceeding with analysis... done in 0m0.10s
```

```
done (passed successfully)
```

# No underspecs, thanks

- Let's make Dialyzer a bit more pedantic

```
$ erlc +debug_info warmup.erl
```

```
$ dialyzer -Wunderspecs warmup.beam
```

```
Checking whether the PLT ... is up-to-date... yes
```

```
Proceeding with analysis...
```

```
warmup.erl:8:2: Type specification warmup:move
```

```
    (any(), any()) -> any() is a supertype of  
      the success typing: warmup:move
```

```
    ('north' | 'west', {_, _}) -> {_, _}
```

```
done in 0m0.14s
```

# Let's try to TypeEr

- ▶ TypeEr is a Type-annotator for Erlang programs
- ▶ Let's see how well it can type move/2:

```
$ typer warmup.erl
%% File: "warmup.erl"
%% -----
-spec move('north' | 'west' ,{_,_}) -> {_,_}.
```

- ▶ 'north' | 'west' is definitely more specific than any()!
- ▶ {\_,\_} is a bit more specific than any(), but not terribly so
- ▶ Type inference in Erlang/OTP, in general, is rather weak
  - ▶ Using '+' does not induce numeric types for X and Y

# Singleton and Predefined Types

- ▶ A singleton type is an Erlang term
  - ▶ For instance, true, false, 1, 42, "Homer",  
{ } (empty tuple), [ ] (empty list), #{} (empty map)
- ▶ Predefined types represent infinite sets of terms

Type	Meaning
any()	Any value
atom()	Any atom
integer()	Any integer <sup>4</sup>
tuple()	Any tuple
float()	Any float
fun()	Any function
pid()	Any process identifier

---

<sup>4</sup>Erlang uses arbitrary-precision arithmetic

## Composite Types

- Composed using, among others, the following forms:

Syntax	Name	Example
$T_1 \mid \dots \mid T_n$	Union	<code>integer()   float()</code>
<code>[ T ]</code>	List	<code>[integer()]</code>
<code>{ T<sub>1</sub> , ... , T<sub>n</sub> }</code>	Tuple	<code>{integer(), integer()}</code>
<code>#{ T<sub>1</sub> =&gt; T'<sub>1</sub> , ... }</code>	Map	<code>#{atom() =&gt; integer()}</code>

- Some predefined composite types:

Type	Defined as
<code>boolean()</code>	<code>true   false</code>
<code>list()</code>	<code>[any()]</code>
<code>string()</code>	<code>[char()]</code>
<code>map()</code>	<code>#{any() =&gt; any()}</code>

- For more, see Erlang/OTP documentation<sup>5</sup>

<sup>5</sup>[http://erlang.org/documentation/doc-12.0/doc/reference\\_manual/typespec.html#builtin\\_types](http://erlang.org/documentation/doc-12.0/doc/reference_manual/typespec.html#builtin_types)

# Function Types

$\text{'fun(' '(' } T_1^{param} \text{' , } \dots \text{' , } T_m^{param} \text{' )' ' -> ' } T^{res} \text{' )'}$

- Consider the function literal

```
Move = fun(Dir, {X,Y}) ->  
    case Dir of  
        north -> {X, Y+1};  
        west  -> {X-1, Y}  
    end  
end.
```

- Valid typings include:

- $\text{fun(any(), any()) -> any()}$
- $\text{fun(atom(), tuple()) -> tuple()}$
- $\text{fun(atom(), \{integer(), integer()\})}$   
 $\text{-> \{integer(), integer()\}}$



# Typespecs

- A substantially better typing of move/2:

```
-module(warmup).
```

```
-export([move/2]).
```

```
-type position() :: {integer(), integer()}.
```

```
-type direction() :: north | west | east | south.
```

```
-spec move(direction(), position()) -> position().
```

```
move(north, {X, Y}) -> {X, Y+1};
```

```
move(west, {X, Y}) -> {X-1, Y}.
```

## Part IV

### Less Basic Concepts

# Exceptions

- We can catch exceptions using **try**:

```
try Expr of
  Pat1 -> Expr1;
  Pat2 -> Expr2;
  ...
catch
  ExPat1 -> ExExpr1;
  ExPat2 -> ExExpr2;
  ...
after
  AfterExpr
end
```

- And we can throw an exception using **throw**:  
throw(Expr)

## Exceptional Moves

```
-module(exceptional_moves).  
-export([move/2,ignore_invalid/2]).
```

```
move(north, {X, Y}) -> {X, Y+1};  
move(west, {0, _}) -> throw(invalid_move);  
move(west, {X, Y}) -> {X-1, Y}.
```

```
ignore_invalid(Dir, Pos) ->  
    try move(Dir, Pos)  
    catch  
        invalid_move -> Pos  
    end.
```

# Algebraic Data Types, or lack thereof

- ▶ In Erlang, we use tuples and atoms to build data structures.
- ▶ Representing trees in Haskell

```
data Tree a = Leaf | Node a (Tree a) (Tree a)
t = Node 6 (Node 3 Leaf Leaf) (Node 9 Leaf Leaf)
```

- ▶ Representing trees in Erlang

```
T = {node, 6, {node, 3, leaf, leaf},
      {node, 9, leaf, leaf}}.
```

# Traversing Trees

- ▶ in Haskell:

```
contains _ Leaf = False
contains key (Node k left right) =
    if key == k then True
    else if key < k then contains key left
         else contains key right
```

- ▶ in Erlang:

```
contains(_, leaf) -> false;
contains(Key, {node, K, Left, Right}) ->
    if Key == K -> true;
       Key < K   -> contains(Key, Left);
       Key > K   -> contains(Key, Right)
    end.
```

# Binary Data

- ▶ Erlang have outstanding support for working with raw byte-aligned data (**binaries**)
- ▶  $\langle\langle b_1, b_2, \dots, b_n \rangle\rangle$  is an  $n$ -byte value
  - ▶ 8-bit:  $\langle\langle 111 \rangle\rangle$
  - ▶ 32-bit:  $\langle\langle 0, 0, 0, 0 \rangle\rangle$
  - ▶ 40-bit:  $\langle\langle \text{"Homer"} \rangle\rangle$
- ▶ Bit Syntax is used to pack and unpack binaries, here we can specify the size and encoding details (like endianness, for instance) for each element of the binary
  - ▶ General form:

$$\langle\langle E_1, E_2, \dots, E_n \rangle\rangle$$

where each element  $E$  have the form:

$$V : \text{size/type}$$

where  $V$  is a value and *size* and *type* can be omitted.

## 8-Bit Colour

- Suppose we need to work with 8-bit colour images, encoded in RGB format with 3 bits for the red and green components and 2 bits for the blue component.



## 8-Bit Colour

- Suppose we need to work with 8-bit colour images, encoded in RGB format with 3 bits for the red and green components and 2 bits for the blue component.
- Pack and unpack functions:

```
pack8bit(R,G,B) ->  
  <<R:3,G:3,B:2>>.
```

```
unpack8bit(<<R:3,G:3,B:2>>) ->  
  {R, G, B}.
```

## Part V

# Concurrency Primitives

# Concurrency-Oriented Programming

- ▶ The world is concurrent
  - ▶ Many different things happen at the same time
- ▶ When we write programs that model, or interact with the world, concurrency should be easy to model

# Parallelism $\neq$ Concurrency

## ▶ Parallelism

- ▶ Bursts of non-interruptible computation, running on multiple processing units
- ▶ Fixed synchronization points
- ▶ Examples: GPUs, TPUs
- ▶ Maximise amount of computation per clock cycle

## ▶ Concurrency

- ▶ Interleaving threads of execution, running on one, or multiple processing units
- ▶ Sporadic communication
- ▶ Suitable for modeling, and interacting with the world
- ▶ Minimise latency

# Concurrency In Erlang

- ▶ Structure system in terms of lightweight, independent processes
- ▶ Processes can only communicate through message passing
- ▶ Message passing is fast
- ▶ Message passing is asynchronous (mailbox model)

# Processes

- ▶ Processes can only communicate through message passing
- ▶ All processes have a unique process ID (pid)
- ▶ Any value can be sent and received (serialization)

# Processes

- ▶ Processes can only communicate through message passing
- ▶ All processes have a unique process ID (pid)
- ▶ Any value can be sent and received (serialization)
- ▶ We can **send** messages:

Pid ! Message

(we can get our own pid by using the build-in function `self/0`)

## Receiving messages

- ▶ Mailbox ordered by arrival – *not* send time



# Receiving messages

- Mailbox ordered by arrival – *not* send time
- We can **receive** messages:

```
receive  
  Pat1 -> Expr1;  
  Pat2 when ... -> Expr2;  
  ...  
after  
  Time -> TimeOutExpr  
end
```

times-out after `Time` milliseconds if we haven't received a message matching one of `Pat1`, `Pat2` with side condition, ....

- **receive** ... **end** is an expression (just like **case** and **if**).

# Spawning Processes

- We can `spawn` new processes:

```
Pid = spawn(Fun)
```

or

```
Pid = spawn(Module, Fun, Args)
```

# Concurrency Primitives, Summary

- ▶ We can **spawn** new processes:

```
Pid = spawn(Fun)
```

(we can get our own pid by using the build-in function `self`)

- ▶ We can **send** messages:

```
Pid ! Message
```

- ▶ We can **receive** messages:

```
receive
```

```
  Pat1 -> Expr1;
```

```
  Pat2 -> Expr2;
```

```
  ...
```


```
after
```

```
  Time -> TimeOutExpr
```

```
end
```


where we get a time-out after `Time` milliseconds if we haven't received a message matching one of `Pat1`, `Pat2`, ....

## Student Activation: Define your favourite function



Make a module, fib, with your favourite (recursive) function.

## Student Activation: Make It Concurrent



## Part VI

### State – How To Deal With It

# Dealing With State

- ▶ Functions are pure (stateless).
- ▶ Processes are stateful.
- ▶ We organise our code as micro-servers that manage a state which can be manipulated via a client API.

## Client-Server Basic Set Up

- ▶ We often want computations to be made in a server process rather than just in a function.
- ▶ That is, we start with something like the following template:

```
start() -> spawn(fun loop/0).  
request_reply(Pid, Request) ->  
    Pid ! {self(), Request},  
    receive  
        {Pid, Response} -> Response  
    end.  
loop() ->  
    receive  
        {From, Request} ->  
            From ! {self(), ComputeResult(Request)},  
            loop();  
        {From, OtherReq} ->  
            From ! {self(), ComputeOther(OtherReq)},  
            loop()  
    end.
```



## Example: Position Server

- ▶ Make a server that can keep track of a *position*.
- ▶ We should be able to:
  - ▶ move the position in some direction
  - ▶ get\_pos to get the position

## Example: Position Server, Implementation

```
start(Start) -> spawn(fun () -> loop(Start) end).  
move(Pid, Dir) -> request_reply(Pid, {move, Dir}).  
get_pos(Pid) -> request_reply(Pid, get_pos).
```

```
request_reply(Pid, Request) ->  
  Pid ! {self(), Request},  
  receive  
    {Pid, Response} -> Response  
  end.
```

```
loop({X,Y}) ->  
  receive  
    {From, {move, north}} ->  
      From ! {self(), ok},  
      loop({X, Y+1});  
    {From, {move, west}} ->  
      From ! {self(), ok},  
      loop({X-1, Y});  
    {From, get_pos} ->  
      From ! {self(), {X,Y}},  
      loop({X,Y})  
  end.
```

## Student Activation: Count Server

- ▶ Let's make a server that can keep track of a counter
- ▶ What is the client API?
- ▶ What is the internal state?

## Example: Todo-List, Interface

```
start() -> spawn(fun() -> loop([]) end).
```

```
add_item(Pid, Description, Due) ->  
    request_reply(Pid, {add, {Description, Due}}).
```

```
all_items(Pid) ->  
    request_reply(Pid, all_items).
```

```
finish(Pid, Index) ->  
    request_reply(Pid, {finish, Index}).
```

## Example: Todo-List, Internal loop

```
loop(Items) ->
  receive
    {From, {add, {Description, Due}}} ->
      Item = #{ description => Description, due => Due},
      From ! {self(), ok},
      loop([Item | Items]);
    {From, all_items} ->
      From ! {self(), {ok, Items}},
      loop(Items);
    {From, {finish, Index}} ->
      Len = length(Items),
      if Index <= 0; Len < Index ->
        From ! {self(), {error, index_out_of_bounds}},
        loop(Items);
      Index > 0, Index <= Len ->
        {L1, [_ | L2]} = lists:split(Index-1, Items),
        From ! {self(), ok},
        loop(L1++L2)
      end
  end.
end.
```

## Part VII

### Summary

## Common Erlang Pitfalls

- ▶ Variables starts with an upper-case letter, atoms starts with a lower-case letter.
- ▶ Erlang does not have statements, only expressions.
- ▶ if expressions (you need to understand what a *guard expression* is).
- ▶ Misunderstanding how patterns works.
- ▶ Functions starts processes, processes runs functions, functions are defined in modules.
- ▶ Not realising when to use asynchronous communication and when to use synchronous communication.

# Summary

- ▶ Parallelism is not the same as concurrency
- ▶ Share-nothing (that is, immutable data) and message passing takes a lot of the pain out of concurrent programming
- ▶ Erlang code is hard to type check, so test it!
- ▶ Keep your mailboxes tidy